



Ask a paper. Check every answer.

Chat with one paper at a time. Answers cite the passages they used, so you can check each claim on the page.


What method do they use?

BERT masks 15% of its input tokens and learns to predict them [C1].



C1

C2



Swipe →



1 • ASK

Type a question, or pick a suggested one.

A short paper is sent to the model whole. A long one is searched first for its most relevant passages.



Ask this paper

Answers cite the passages they use. Click a citation to see it in the paper.

Summarize the main contribution

What method do they use?



What are the limitations?

Ask about this paper...





2 • READ THE ANSWER

Answers point to the passages they used.

The answer streams in as it's written, with citations like [C1] that you can hover and click.

What method do they use?

BERT is pre-trained on two tasks. It masks 15% of the input tokens and learns to predict them [C1]. It also learns to tell whether one sentence follows another [C2].



C1

C2

3 • CHECK THE SOURCE


Click a citation. Land on the passage.

Hover [C1] to see its page and section. Click it, and the paper scrolls to that passage and flashes it.

Source 1: page 4, section “3.1 Pre-training BERT”

The AI used this passage. Click to see it in the paper.

It masks 15% of the input tokens
and learns to predict them [\[C1\]](#).



BERT paper · p. 4





4 · KEEP WHAT HELPS

Save part of an answer as a note.

Select it and click Save as note. The note is anchored on the passage it cites and marked AI.

BERT is pre-trained on two tasks. It masks 15% of the input tokens and learns to predict them [C1]. It also learns to tell whether one sentence follows another [C2].

✧ Save as note

✧ AI

p. 4

In order to train a deep bidirectional representation, we simply mask some percentage of the input tokens at random...


It masks 15% of the input tokens and learns to predict them [C1].



PRIVATE BY DEFAULT

Runs on your machine. AI is always labelled.

With the default Ollama model, the text of your papers never leaves your computer. Each answer names its model.





**Read deeper.
Keep every thought.
Ask, then check.**

PaperLab runs on your own machine. Your PDFs, notes and chats stay on your computer.

■ github.com/khannoussi-malek/PaperLab